BATCH 9 · ONE 150-WORD NOTE PER MEMBER

Individual Reflections

Credit Pre-Screening Engine · MUC02 · Path 2, KLU Batch

DRAFT — each member should read their own note and edit it into their own words before submission. These were written from each member's section of the notebook and the decisions recorded there. The reflection is meant to be personal, so change anything that does not sound like you.

Mohammed Mutee Ruknuddin

Data Analyst — Step 5.1 and the CBIM Problem Canvas

I owned the cleaning and the problem canvas. The part I expected to be routine turned out to be the part that mattered most. Three issues needed a decision, and the order mattered: I corrected twelve annual incomes that had been typed into a monthly field first, because leaving them in would have poisoned the employment-type medians I then used to fill 441 missing incomes. I kept the 1,039 blank credit scores after checking that credit history was zero for every one of them — they were new-to-credit applicants, not errors. What I learned from Javeed's section is that my cleaning choices moved his recommended threshold between 0.85 and 0.90, worth about 65 analyst hours a month. A cleaning log is not paperwork; it is the audit trail for a business number.

[132 words]

Mohammed Saalif Udyawar

Insights Analyst — Steps 5.2 and 5.3

I looked for what the credit team actually does before anyone built a model. Two cliffs came out of it. Below a credit score of 650 only 3.0% of applications are approved against 32.5% just above — a rule, not a gradient. And once EMIs pass half of income, approvals collapse from 47.4% to 15.6%. FOIR did not exist in the file; I had to build it, and it turned out to be the strongest feature in the whole dataset. What I learned from Mutee's section is how much a finding depends on the cleaning underneath it — my FOIR cliff only holds because he corrected those twelve incomes first. I also learned from Javeed that a number can be real in a table and still mislead in a model, which is what happened to loan-to-income.

[136 words]

Mulla Javeed

Model Builder — Steps 5.4 and 5.5

I built the logistic regression and chose the threshold. The model reached 79.1% accuracy and 0.867 ROC-AUC, but the number that changed the conversation was 225: at the default 0.5 cut-off, 225 applicants the bank would have approved get turned away. That is what made the threshold table necessary. At 0.90 we auto-decline 15.2% of applications with 4.7% wrong, just inside the

bank's 5% limit — about 456 applications and 190 analyst hours a month. The model also rediscovered both of Saalif's cliffs on its own, which is the best evidence we have that it learned real policy. From Varshitha's section I learned that a model being accurate overall says nothing about whether it is accurate for the people it is about to reject.

[124 words]

Nagothu Varshitha

Policy Synthesiser — Steps 5.6 and 5.7

I ran the fairness check and wrote the briefing. The historical data shows a 3.7-point approval gap between male and female applicants, but it nearly vanishes inside Gig/Contract and Business Owner applicants — the gap comes from which groups women are in, not from different treatment within them. Our model does not widen it: 15.6% of women are auto-declined against 15.0% of men. The harder finding was that leaving gender out does not make a model fair, because income and employment type are strong proxies for it. From Javeed's section I learned that the model is only 56.7% accurate on new-to-credit applicants, which is why my briefing recommends they never enter the auto-decline lane at all.

[116 words]